

A Verified Garbage Collector for Gallina

Shengyi Wang, Anshuman Mohan, Aquinas Hobor



APLAS NIER
November 18, 2019

Verify **graph-manipulating** programs
written in **executable C**
with **machine-checked** correctness proofs

Verify **graph-manipulating** programs
written in **executable C**
with **machine-checked** correctness proofs

Ubiquitous in critical areas



Certifying Graph-Manipulating C Programs via Localizations within Data Structures

SHENGYI WANG, National University of Singapore, Singapore

QINXIANG CAO, Shanghai Jiao Tong University, China

ANSHUMAN MOHAN, National University of Singapore, Singapore

AQUINAS HOBOR, National University of Singapore, Singapore

VST + CompCert + 40000 LOC library



Certifying Graph-Manipulating C Programs via Localizations within Data Structures

SHENGYI WANG, National University of Singapore, Singapore

QINXIANG CAO, Shanghai Jiao Tong University, China

ANSHUMAN MOHAN, National University of Singapore, Singapore

AQUINAS HOBOR, National University of Singapore, Singapore

VST + CompCert + 40000 LOC library

Powerful enough to verify **executable code**
against **realistic specifications**
expressed with **mathematical graphs**



Certifying Graph-Manipulating C Programs via Localizations within Data Structures

SHENGYI WANG, National University of Singapore, Singapore

QINXIANG CAO, Shanghai Jiao Tong University, China

ANSHUMAN MOHAN, National University of Singapore, Singapore

AQUINAS HOBOR, National University of Singapore, Singapore

VST + CompCert + 40000 LOC library

Powerful enough to verify **executable code**
against **realistic specifications**
expressed with **mathematical graphs**

[Wang *et. al.*, PACMPL OOPSLA 2019]



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly

This Talk



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly

Gallina assumes **infinite** memory
but CompCert C has a **finite** heap



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly

Gallina assumes **infinite** memory
but CompCert C has a **finite** heap

Solution: garbage collect the CompCert C code



Gallina \rightsquigarrow CompCert C \rightsquigarrow Assembly

Gallina assumes **infinite** memory
but CompCert C has a **finite** heap

Solution: garbage collect the CompCert C code

New problem: verify the garbage collector

GC has jurisdiction over the heap




GC has jurisdiction over the heap



Garbage Collection: an Introduction

GC has jurisdiction over the heap

Mutator mallocs in special subheap




Garbage Collection: an Introduction

GC has jurisdiction over the heap

Mutator `mallocs` in special subheap

If subheap is full




Garbage Collection: an Introduction

GC has jurisdiction over the heap

Mutator `mallocs` in special subheap

If subheap is full **call GC**




Garbage Collection: an Introduction

GC has jurisdiction over the heap


Mutator `mallocs` in special subheap

If subheap is full **call GC** and try again




Primum non nocere: first, do no harm


Primum non nocere: first, do no harm



Primum non nocere: first, do no harm



Primum non nocere: first, do no harm




- 12 generations, doubling in size
- Functional mutator: no back pointers


- 12 generations, doubling in size
- Functional mutator: no back pointers
- Cheney's mark-and-copy collects gen to next
- Potentially triggers cascade of pairwise collections

- 12 generations, doubling in size
- Functional mutator: no back pointers
- Cheney's mark-and-copy collects gen to next
- Potentially triggers cascade of pairwise collections
- Three key functions:
 - `forward` copies individual objects
 - `do_scan` repairs copied objects
 - `forward_roots` kick-starts the collection

Overview of Operation'




Overview of Operation'




forward ✓

Overview of Operation'



forward ✓


Overview of Operation'



forward ✓

do_scan ✓

Overview of Operation'




forward ✓

do_scan ✓

do_gen ✓


Non-Concerns

roots
(2,3)
(2,2)
(3,1)



more garbage? ✗

Non-Concerns




more garbage? ✗

back pointers? ✗


Sources of Complexity

variable-length objects



Sources of Complexity

variable-length objects
on-the-fly int/ptr disambiguation



forward is robust

forward is robust
pointer?

forward is robust
pointer?
in from space?

forward is robust
pointer?
in from space?
already forwarded?

forward is **robust** and **versatile**
pointer?
in from space?
already forwarded?

forward is **robust** and **versatile**

pointer? called on root set

in from space?

already forwarded?

forward is **robust** and **versatile**

pointer? called on root set

in from space? called on heap

already forwarded?

forward is **robust** and **versatile**

pointer? called on root set

in from space? called on heap

already forwarded?

Specifying forward *functionally* is too hard

forward is **robust** and **versatile**

pointer?	called on root set
in from space?	called on heap
already forwarded?	

Specifying forward *functionally* is too hard

forward_relation explains how
the graph may change because of forward

```
Inductive forward_relation (from to : nat) :  
  forward_t -> LGraph -> LGraph -> Prop :=
```

```
Inductive forward_relation (from to : nat) :  
  forward_t -> LGraph -> LGraph -> Prop :=  
| fr_v_not_in : forall (v : VType) (g : LGraph),  
  vgeneration v <> from ->  
  forward_relation from to (inl (inr v)) g g
```

```
Inductive forward_relation (from to : nat) :  
  forward_t -> LGraph -> LGraph -> Prop :=  
| fr_v_not_in : forall (v : VType) (g : LGraph),  
  vgeneration v <> from ->  
  forward_relation from to (inl (inr v)) g g  
| fr_e_to_fwdded : forall (e : EType) (g : LGraph),  
  vgeneration (dst g e) = from ->  
  raw_mark (vlabel g (dst g e)) = true ->  
  let new_g := labeledgraph_gen_dst g e  
    (copied_vertex (vlabel g (dst g e))) in  
  forward_relation from to (inr e) g new_g
```

```
| fr_e_to_not_fwdded_Sn : forall (e : EType) (g g' : LGraph),  
  vgeneration (dst g e) = from ->  
  raw_mark (vlabel g (dst g e)) = false ->  
  let new_g :=  
    labeledgraph_gen_dst (lgraph_copy1v g (dst g e) to)  
    e (copy1v_new_v g to) in forward_loop from to  
    (make_fields new_g (copy1v_new_v g to)) new_g g' ->  
    forward_relation from to (inr e) g g'
```

```

| fr_e_to_not_fwdded_Sn : forall (e : EType) (g g' : LGraph),
  vgeneration (dst g e) = from ->
  raw_mark (vlabel g (dst g e)) = false ->
  let new_g :=
    labeledgraph_gen_dst (lgraph_copy1v g (dst g e) to)
    e (copy1v_new_v g to) in forward_loop from to
    (make_fields new_g (copy1v_new_v g to)) new_g g' ->
    forward_relation from to (inr e) g g'

with forward_loop (from to : nat) :
  list field_t -> LGraph -> LGraph -> Prop :=
| fl_nil : forall (g : LGraph), forward_loop from to [] g g
| fl_cons : forall (g1 g2 g3 : LGraph)
  (f : field_t) (fl : list field_t),
  forward_relation from to (field2forward f) g1 g2 ->
  forward_loop from to fl g2 g3 ->
  forward_loop from to (f :: fl) g1 g3

```

Similar to `forward_relation`, we have

`forward_roots_relation`

`do_scan_relation`

`do_generation_relation`

`garbage_collect_relation`

Similar to `forward_relation`, we have

`forward_roots_relation`

`do_scan_relation`

`do_generation_relation`

`garbage_collect_relation`

A composition of these gives us our **isomorphism**

More Meat and Potatoes

- Cheney implemented too conservatively:
only part of `to` space needs to be scanned

- Cheney implemented too conservatively:
only part of `to` space needs to be scanned
Performance doubled

- Cheney implemented too conservatively:
only part of `to` space needs to be scanned
Performance doubled

- Overflow in the following calculation:

```
int space_size =  
    h->spaces[i].limit - h->spaces[i].start;
```

- Cheney implemented too conservatively:
only part of `to` space needs to be scanned
Performance doubled


- Overflow in the following calculation:

```
int space_size =  
    h->spaces[i].limit - h->spaces[i].start;
```

Fixed by adjusting nursery size


Double-bounded pointer comparisons:

```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi); }
```




Double-bounded pointer comparisons:

```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi); }
```




Double-bounded pointer comparisons:

```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi);  
}
```



Double-bounded pointer comparisons:


```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi); }
```



Undefined behavior in C

Double-bounded pointer comparisons:

```
int Is_from(value * lo, value * hi, value * v) {  
    return (lo <= v && v < hi); }  
}
```



Resolved using CompCert's `extcall_properties`

A classic OCaml trick:

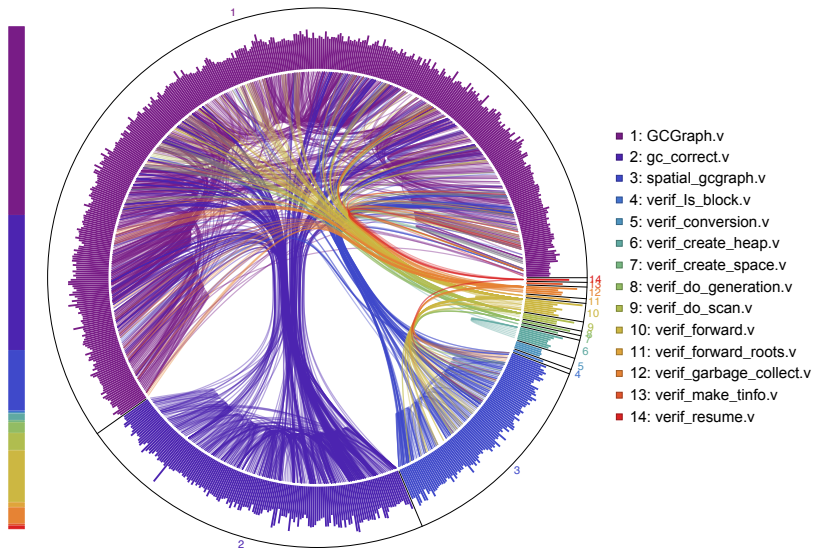
```
int test_int_or_ptr (value x) {  
    return (int)(((intnat)x)&1); }
```

A classic OCaml trick:

```
int test_int_or_ptr (value x) {  
    return (int)(((intnat)x)&1); }
```

Discussing char alignment issues with CompCert

Reusability: separation between pure and spatial reasoning



Problems of a similar shape

Problems of a **similar shape**
serialization
other collectors

Problems of a similar shape
 serialization
 other collectors

Towards a verified GC for OCaml

Problems of a **similar shape**
 serialization
 other collectors

Towards a verified GC for **OCaml**
 mutability
 calculate root set
 allow other datatypes

Problems of a **similar shape**
serialization
other collectors

Towards a verified GC for **OCaml**
mutability
calculate root set
allow other datatypes

Further refinements required in C semantics
before we can **specify** and **verify** OCaml's GC?